

Advanced Machine Learning for Physics student project: RECENTRE

Francesco Il Passero and Eugenio Nicotina
(Dated: July 10, 2026)

We build a machine learning framework to train, validate, finetune models and perform data analysis for the RECENTRE (REal-time motion CorrEction in magNeTic REsonance) project. We first reproduce the original results of the RECENTRE project based on a GRU network through a config-oriented software design, then we introduce and evaluate more modern machine learning architectures, studying in what regimes they are most effective and whether they're better candidates for real time deployment in MRI machines. Model size, inference latency, robustness to noise and uncertainty calibration are among the many different metrics by which we evaluate the performances of each of the proposed architectures.

I. INTRODUCTION

The RECENTRE project is concerned with real time prediction and correction of head motion parameters. A crucial fact to take into account when building models for this task is the inference latency of the model. The temporal distance between the frame recording and the subsequent next frame prediction must be of the order of 100 ms, so the inference latency must be well below this limit to account for other software- and hardware-related latencies.

The main goals of the present project are to reproduce the original paper's ([1]) results, which are based on a GRU network, and evaluate modern machine learning architectures such as transformers, TCNs, and physics-informed models to study whether they are better candidates for the project's task.

Given the large amount of models and the large amount of training runs (different model hyperparameters, different train / test tasks, data augmentation, etc...) we introduced a config-based software design pattern to keep the code and the output data well organized. This way, new models can be implemented and registered in a model registry dictionary; the hyperparameters of the model, the preprocessing pipeline of the data and the training details (epochs, learning rate, ...) can all be specified in a minimal, easy-to-read configuration file. This way, we built a solid end-to-end pipeline from the original raw data and model specification to the full training and evaluation loop.

II. DATASET

The raw dataset is taken from the Human Connectome Project (HCP), which consists in navigator data taken by MRI machines with a sampling frequency of 720 ms. Each 'frame' has 6 motion parameters: 3 of them are translation parameters describing the head position of the patient, the remaining 3 are rotation parameters indicating the inclination of the patient's head. The dataset contains three different scenarios: the "R" (resting) one, where patients were resting while the MRI scan was running, then the "M" (memory) one, where patients

were recollecting memories, and the "L" (language) one, where patients were actively talking. There are a total of 3221 time series among the three tasks. Each time series is associated with a patient. The preprocessing pipeline discards 140 time series, so that 1027 patients are all associated with the 3 time series of the "R", "M", "L" tasks. The three tasks have different time lengths. The R data is 1200 frames long, while the M and L are 405 and 316 frames long, respectively. Lots of efforts were put in the thesis and paper ([2]) to study different data splits for training and testing, e.g. considering the "cross-task" scenario or "cross-patient" scenario. We build on their findings (no substantial difference in performance among the different scenarios) by:

1. Developing a minimal, simple pipeline to choose training and testing tasks and choosing whether to use different splits of patients
2. Standardizing all results to the same scenario: train on all tasks (R+M+L) and test on all tasks (R+M+L), while keeping disjoint sets of patients to avoid any possible leakage




FIG. 1: Example patient time series showing predicted vs. true head-motion parameters across the six degrees of freedom.

A subsequent analysis on physics-informed models has shown that exploiting additional kinematic data such as velocity and acceleration parameters substantially im-

proved test performances. So, 6 additional columns containing velocity parameters and 6 columns containing acceleration parameters for the 6 different dimensions are kept and used.

A. Data augmentation

The dataset contains roughly 2 million input window - frame target pairs, so data augmentation is not a priority for the project. Nevertheless, we investigated whether data augmentation could provide a benefit to model performances. Since the translation and rotation parameters are physically linked, there is no much room for independent transformations of the input data. However, two possible transformations are considered for data augmentation: time reversal and sign flip. These two augmentation techniques can quadruple the initial training dataset.

B. Preprocessing

The preprocessing pipeline consists in extracting the raw data from the HCP files, converting degrees to radians, ensuring a standardized time length for each of the different tasks, and ensuring that each patient is associated with its three runs. The preprocessing pipeline is meant to be run only once, as it builds three processed dictionaries (one for each task) that associate each patient id with its time series for that task. All successive scripts read from the processed datasets.

III. METHODOLOGY

A. Pipeline

All the tested models follow the same protocol:

1. Model definition and registration
2. Training config compilation
3. Training loop
4. Evaluation loop and data analysis
5. Optional fine-tuning

After model definition, typically as a torch module, a training config is necessary to run the training loop. The model hyperparameters, training and testing tasks, split percentages, data augmentation, loss function, training epochs, learning rate and so on are all specified in the training config, so that all the relevant parameters are easy to read, especially for systematic comparisons between many different models and runs.

After the training is done, the model weights, the optimizer and scheduler states, training configs, info about

the training and testing tasks and patient ids are stored in a checkpoint to be used for testing, finetuning or resuming training.

B. Metrics

The fundamental metric that describes the prediction error of a single frame is called frame displacement (FD). It is defined as

$$FD(t) = \sum_{i=1}^3 |y_i(t) - \hat{y}_i(t)| + 50\text{mm} \sum_{i=4}^6 |y_i(t) - \hat{y}_i(t)| \quad (1)$$

it is essentially an absolute error that combines translation error ($i = 1, 2, 3$) and rotation error ($i = 4, 5, 6$), where the radians are scaled by 50mm, the average skull radius. This metric is a single scalar describing the absolute error of the prediction relative to the true position. Since the sampling rate is 720ms, the lag1-model that predicts the last observed position as the next one is already a solid baseline. We'll refer to the lag-1 model as the baseline model. Naturally, we can define the FD_{base} as the frame displacement for the lag1-model:

$$FD_{base}(t) = \sum_{i=1}^3 |y_i(t) - y_i(t-1)| + 50\text{mm} \sum_{i=4}^6 |y_i(t) - y_i(t-1)| \quad (2)$$

Another useful metric is the FD_{gain} , that is the relative difference between FD_{base} and FD_{model} :

$$FD_{gain} = \frac{FD_{base} - FD_{model}}{FD_{base}} \quad (3)$$

We aim for large, positive FD_{gain} .

Another useful metric is the percentage of frames that is improved (i.e. lower frame displacement) with respect to the baseline. Analogously, we can look at the percentage of patients with an average $FD_{gain} > 0$.

C. Loss function

We use the same form of the loss function as the original work:

$$L = L_{NLL} - \beta FD_{gain} \quad (4)$$

Where L_{NLL} is the gaussian negative log likelihood term that is needed for both prediction accuracy and uncertainty calibration.

Using FD_{gain} in the loss is a way to penalize the model if it does worse than the baseline model, especially when the baseline model is performing well in that particular frame. This loss term is essentially specifying that we'd rather accept large FDs in frames where the baseline has poor performance than have frames where the model is worse than the baseline.

We also reduce the learning rate on loss plateau and use early stopping.

D. Finetuning

Finetuning in this project is intended as a per-task, per-patient finetuning. After a model has been pre-trained on many different time series, it can be additionally fine-tuned on the first part of a time series so that the model can be adapted to the particularities of the patient and task under observation. The original work has found a substantial improvement in test performance after finetuning, relative to the pre-trained generalist model.

The finetuning implementation idea is the same as the original work one. It uses a modified loss function:

$$L = L_{MSE} - \beta FD_{gain} + \lambda L_{L2SP} \quad (5)$$

It differs from the pre-training loss function because it lacks an uncertainty calibration term (which is fine because the uncertainty head is frozen) and it has a L^2 starting point (L2SP) term that penalizes significant drifts from the pre-trained model:

$$L_{L2SP} = \sum_i (w_i - w_i^0)^2 \quad (6)$$

where w_i^0 is the i -th pretrained model weight.

Our own implementation aims to standardize this procedure across all the implemented models, and make all the finetuning parameters readily accessible in a dedicated finetuning config file. Each model inherits from a TunableModel base class that automatically splits the model parameters into a frozen variance head, a higher learning rate mean head and a lower learning rate body. The learning rates and the dropout rate are configurable in the finetuning config.

IV. REPRODUCTION OF PREVIOUS RESULTS

The first part of the project is concerned with reproducing the results of the previous thesis and paper. First, we refactored the codebase, fixed minor bugs and organized it in a config-oriented design, with standardization of train/test scenarios, model definitions, finetuning pipeline, model / data / train hyperparameters, data analysis and visualizations.

Then, we focused on reproducing these GRU results from the previous work:

1. Paper main result: $R^2 > 0.87$ for all the 3 tasks and 6 dimensions.
2. Thesis result: roughly 70% frames have $FD_{gain} > 0$.
3. Thesis result: positive FD_{gain} frames percentage rises to 75% for the finetuned model.

We first reproduced the original results by considering the per-task training and testing, but we soon decided to

directly improve these results by adopting some fairly obvious refinements on top of the original structure. First, we dropped the cross-task scenario debate and decided to train and test a single model on all three tasks. We obviously used different patients for the train, validation and test splits to avoid any data leakage. The hyperparameters of the model are left the same. The results are visible in Fig. (2)

We also reproduced the per-patient fine-tuning result of the previous work. Fig. (3) summarizes the effect of fine-tuning, and Fig. (4) compares each patient before and after fine-tuning.

V. RESULTS

A. β scan

The previous work used a fixed (arbitrary) choice of $\beta = 0.1$. We investigated the performances of the original GRU model as β is swept over many orders of magnitude. In particular, we focused on the most relevant metrics that could sharply depend on the relative importance between the negative log likelihood and the FD_{gain} term. Since the negative log likelihood term is the only one explicitly constraining the uncertainty of the prediction, we expect that as β grows, the uncertainty calibration will get worse. On the other hand, FD_{gain} and the percentage of improved frames (relative to baseline) should rise with β . Our empirical results are summarized in Fig. (5).

As expected, uncertainty calibration degrades as β grows near and past $\beta = 1$, while FD_{gain} rises and then decreases again when FD_{gain} is over-weighted in the loss. This is probably an overfitting phenomenon specific to the FD_{gain} term: even though the loss function at very high β is essentially a maximization of FD_{gain} , this doesn't generalize well to the validation/test set. Across different scans, the three values $\beta = 0.1, 0.5, 1$ showed the best performances. We chose to use $\beta = 0.5$ as a middle ground between the optimally calibrated but lower FD_{gain} focused $\beta = 0.1$ and the less calibrated but higher FD_{gain} focused $\beta = 1$.

TABLE I: Effect of the jerk-penalty weight γ on test-set performance (GRU, sequence length 32).

γ	Mean FD	Mean FD_{gain}	% Frames > base
0	0.1268	0.1778	78.4
0.001	0.1267	0.1785	78.5
0.01	0.1267	0.1767	78.2
0.1	0.1280	0.1786	79.0
0.3	0.1273	0.1798	78.8
0.5	0.1285	0.1769	78.7
0.7	0.1312	0.1693	78.5
1	0.1363	0.1514	77.3

B. Data augmentation

Next, we investigated whether data augmentation could be helpful for this project. Since the translation and rotation parameters are physically related, there is little room for independent transformation of the motion parameters. We investigated two physically plausible transformations: time reversal and negation (sign flip).

TABLE II: Effect of data augmentation on test-set performance.

Augmentation	FD_{pred}	FD_{gain}	% Frames > base
None	0.143	0.097	73.0
Time-reversal	0.146	0.078	70.9
Negation	0.141	0.114	75.2
Both	0.145	0.085	71.9

From the empirical data in Table II we find that, while time reversal augmentation degrades the performances on the (non augmented) test set, negation significantly improves them. From now on, we always train the models with negation-augmented data.

C. γ scan (jerk penalty)

Since the task of interest is a physical motion prediction task, we investigated whether enforcing a kinematic prior could help the model find a better solution. Indeed, injecting symmetry priors or inductive biases into ML models is widely known to improve performance and/or efficiency in many applications (that's what geometrical deep learning is all about), e.g. enforcing SE(3) equivariance in mesh classification or gauge equivariance in lattice gauge theory models. For tasks that do not have an exact symmetry to enforce, an alternative approach is to inject the physical bias in the loss function. In our case, we studied whether a jerk penalty term in the loss could help the model explore more physically plausible solutions and thus lead to better optimization and generalization. We chose to use a jerk penalty because our data is rapidly fluctuating, so using velocity- or acceleration-related penalties probably wouldn't help. Jerk penalty is a bit more theoretically grounded, as it captures the muscles' ability to produce smooth changes in applied force, while still penalizing very rapid changes in movement.

We added the term

$$L_{jerk} = \gamma J(t) \quad (7)$$

to the loss function, where the jerk $J(t)$ is computed via finite differences as:

$$J(t) = a(t) - a(t-2) = x(t) - 3x(t-2) + 3x(t-4) - x(t-6) \quad (8)$$

Where the 2-strided input is kept to avoid using in-between frames that the model never sees. We then swept γ over many orders of magnitude to study its effect on test performances for the GRU model. The results are shown in Table I.

We see from the results that we gain at most 0.002 ($\sim 1\%$) from the base model. This very small gain is not sufficient to conclude that the jerk penalty improves prediction accuracy. A possible future extension of this project could be to train many models and to associate a statistical error to the test performance so that an actual statistical edge could be confirmed or ruled out.

D. NLinear

The lightest model we implemented is the NLinear model [1] with $\mathcal{O}(L)$ parameters, L being the sequence length. The architecture philosophy is to center the whole window with respect to the last frame and apply six independent linear layers (one for each channel) to predict the shift of the next frame, and another six to predict the associated variance.

E. DLinear

The DLinear model [1] is among the lightest models as well. It separates the global movement (*trend*) with respect to the high frequency one (*season*) by shifting a kernel along the window and taking the mean each kernel-size frames. Then six linear layers each are applied separately to the trend, season and the initial data (for a total of 18 layers). The first two outputs are summed to predict the shift of the next frame while the latter to predict the variances.

F. TSMixer

The Time-Series Mixer [2] (TSMixer) is composed mainly of linear layers as well with the advantage of mixing operations along both the time and feature dimensions. Four TSMixer layers are applied in sequence, each of them composed of a linear layer among the frames and a multi-layer perceptron among the features; plus the activation functions, dropout and residual connections.

G. TCN

Shifting from only linear layers to convolutional ones, the first new model we implemented is a temporal convolutional network [3] (TCN). This network uses temporal convolution kernels at various level of dilation to analyze both long and short range patterns. A schematic diagram of the network is found in Fig. (6).

The TCN model showed remarkable performances, competitive with the original GRU architecture.

H. Transformer

The next model implemented is a basic transformer with two encoder-only layers made of a multi-head self attention across the time frames. The last frame of the output of the transformer is fed to two linear layers that predict the shift of the next frame and the corresponding variance.

I. PatchTST

A transformer-based model for time forecasting is the patch time series Transformer [4] (PatchTST). It enforces a channel-independent analysis: for each channel the window is segmented into multiple patches that are fed as input tokens to the transformer. Also, all channels share the same embedding and transformer weights.

J. Conformer

A conformer model [5] is an architecture that combines transformers and convolutional neural networks (CNN). The former is specialized to analyze long range patterns, while the latter is specialized to analyze short range ones, making the hybrid architecture extremely powerful and indeed the most successful one in our study. From the transformer it inherits the positional embedding, the feed forward and the multi head self-attention. From the CNN it inherits a 1D-convolutional layer with kernel size 1 that doubles the channels number using the same weights across all frames, and performs a Gated Linear Unit (GLU). Another 1D-convolutional layer is applied, this time channel independent with kernel size greater than 1 to capture high frequency patterns, and after the activation function another 1D-convolutional layer with kernel size 1 is applied (followed by dropout) to mix all features before feeding the transformed data to the two linear layers that predict the shift of the next frame and the corresponding variance.

K. Mamba

mamba

L. Routing over experts

Inspired by MoE systems in modern LLMs, we decided to try combining different "experts", i.e. model architectures, to give a better overall frame-by-frame prediction. Given the long training runs of the best (heavy) experts,

we chose to keep the already trained models and just train a global router. Clearly this doesn't let each architecture become an expert in a particular regime, but it is still an interesting ensemble-like strategy. We chose four experts: the best mamba, conformer and GRU models (at sequence length 128), along with the baseline, and trained a router to predict coefficients to linearly combine the prediction of the four models.

The router is a small MLP that takes as input the per-frame experts' predictions and uncertainties and outputs softmax weights over the four models for the current frame. The router is trained to maximize FD_{gain} with the three experts kept frozen. The router is trained on the validation patients and evaluated on the held-out test patients, both unseen by every expert (more precisely, 85% of the validation patients are used for training, 15% are for the actual validation (early stopping) and the test patients are for testing). Results are reported in Table (III).

Model	FD_{gain}
Conformer (single)	0.2159
Mamba (single)	0.2020
GRU (single)	0.1957
Fixed mean	0.2196
Soft router	0.2238
Fixed router-mean weights	0.2221

TABLE III: Routing performances on the test set

We see that the soft routing strategy found by the model does indeed improve performances. We also collected the average weights produced by the router for each model: mamba 16% for the Mamba model, 61% for the Conformer, 15% for the GRU and 8% for the baseline. We immediately see that all models are blended in, and better-performing models are given a larger weight on average. To evaluate the actual role of the per-frame router vs. a fixed linear combination of the models, we evaluated a fixed mean of the 3 best models (33% weight to each model) and a fixed mean of the best models where the weights are given by the previously mentioned averages. We see that active frame-by-frame routing is indeed the best option, although the best fixed linear combination achieves roughly the same performances.

M. Model comparisons

We trained 9 different architectures with 4 sequence lengths each. We compare all of them among very different metrics, from test accuracy to model size to noise robustness. In particular, we rank the models based on FD_{gain} , but we report, for all of them: number of parameters, size, FLOPs latency, throughput, FD, FD_{gain} , $\%FD_{gain} > 0$ NLL, $\Delta FD@0.3\sigma$, tolerance to noise. We report the full results in Table ??.

To have a more visual look at the data, we create a set of Pareto-optimal frontier visualizations. We plot accuracy (FD_{gain}) on the y-axis and cost (latency, FLOPs, size) on the x-axis. We also do a scatterplot of FD_{gain} vs. noise tolerance, although the clear linear term simply shows that all models lose accuracy in the same way when noise is added, there is no clear 'noise-robust' winner, just a better model overall which needs more noise to be brought down in accuracy. The plots are shown in Fig. (??)

N. Distillation

From the pareto-optimal frontier it is clear that the best performing model is the conformer, followed by the mamba and then by the GRU. Given the large difference in latency between the GRU and the conformer, it'd be ideal to push the GRU to have higher accuracy while still keeping their remarkably low latency. To achieve such goal, we tried knowledge distillation from the conformer onto the GRU.


In particular, distillation happens on two separate spaces: the output space (mean, var) and the latent space, just before the fully connected heads. Both mod-

els indeed send the output of their architecture-specific latent space onto a fc mean head and a fc var head. We therefore use the latent space before the fc heads to distill the conformer onto the GRU. Since these latent spaces have different dimensions, we first linearly project the student latent space to the teacher's, and then simply use MSE loss to bring the latent feature closer to each other. For the output space, we use the KL divergence between the normal distributions (mean, var) outputted by the two models. The results are shown in Table ?? and Fig. ?. As we can see, we succeeded in achieving significantly higher FD_{gain} while retaining the very low latency of the GRU model. In future work, this distillation process could be optimized over the various hyperparameters weighing the latent space or output space contributions and the temperature of the teachers' soft labels.


VI. DISCUSSION

VII. CONCLUSIONS


-
- [1] A. Zeng, M. Chen, L. Zhang, and Q. Xu, arXiv preprint arXiv:2205.13504 10.48550/arXiv.2205.13504 (2022).
 - [2] S.-A. Chen, C.-L. Li, N. Yoder, S. O. Arik, and T. Pfister, Tsmixer: An all-mlp architecture for time series forecasting (2023), arXiv:2303.06053 [cs.LG].
 - [3] C. Lea, M. D. Flynn, R. Vidal, A. Reiter, and G. D. Hager, Temporal convolutional networks for action segmentation and detection (2016), arXiv:1611.05267 [cs.CV].
 - [4] Y. Nie, N. H. Nguyen, P. Sinthong, and J. Kalagnanam, A time series is worth 64 words: Long-term forecasting with transformers (2023), arXiv:2211.14730 [cs.LG].
 - [5] A. Gulati, J. Qin, C.-C. Chiu, N. Parmar, Y. Zhang, J. Yu, W. Han, S. Wang, Z. Zhang, Y. Wu, and R. Pang, Conformer: Convolution-augmented transformer for speech recognition (2020), arXiv:2005.08100 [eess.AS].



(a) True vs. predicted value for each of the six motion dimensions, with per-dimension R^2 . All dimensions reach $R^2 > 0.94$.




(b) Framewise-displacement distribution of the model vs. the lag-1 baseline (left) and the per-sample FD_{gain} distribution (right): 74.8% of samples are improved over the baseline.




(c) Per-patient mean FD_{gain} sorted within each task (over 92% of patients improved in every task) and per-patient baseline vs. model FD (rightmost panel).


FIG. 2: Reproduction of the GRU results with a single generalist model trained and tested on all three tasks (R+M+L) over disjoint patient splits. (a) per-dimension agreement, (b) framewise-displacement and per-sample gain distributions, (c) per-patient gains.



(a) Per-patient mean FD_{gain} sorted within each task (over 95% of patients improved in every task) and per-patient baseline vs. model FD (rightmost panel).




(b) Distribution of the per-patient FD_{gain} improvement (fine-tuned - pretrained); the overall mean shift is +0.047.



(c) Per-dimension MSE before and after fine-tuning (rotations scaled by 50^2).

FIG. 3: Per-patient fine-tuning of the generalist GRU model (R+M+L). (a) per-patient gains, (b) distribution of the gain improvement over the pretrained model, (c) per-dimension MSE reduction.




(a) FD_{gain} (win rate 80%).


(b) Predicted FD (points below $y = x$ improved).

(c) Frame-level win rate.

FIG. 4: Per-patient metrics before vs. after fine-tuning, colored by task. Each point is one patient; the dashed line marks no change ($y = x$).



(a) Headline metrics vs. β : test NLL, mean FD vs. baseline, FD_{gain} , and MAE vs. predicted σ .



(b) Calibration vs. β : empirical coverage of the $\pm 1\sigma$ and $\pm 2\sigma$ predictive intervals against their Gaussian targets.

FIG. 5: β scan of the GRU model. As β grows, FD_{gain} peaks around $\beta = 1$ before the loss over-weights the FD term and uncertainty calibration degrades.




FIG. 6: Schematic diagram of the TCN model for a kernel size equal to three, and dilation that doubles at each new layer. The dilation allows far away frames to interact.